

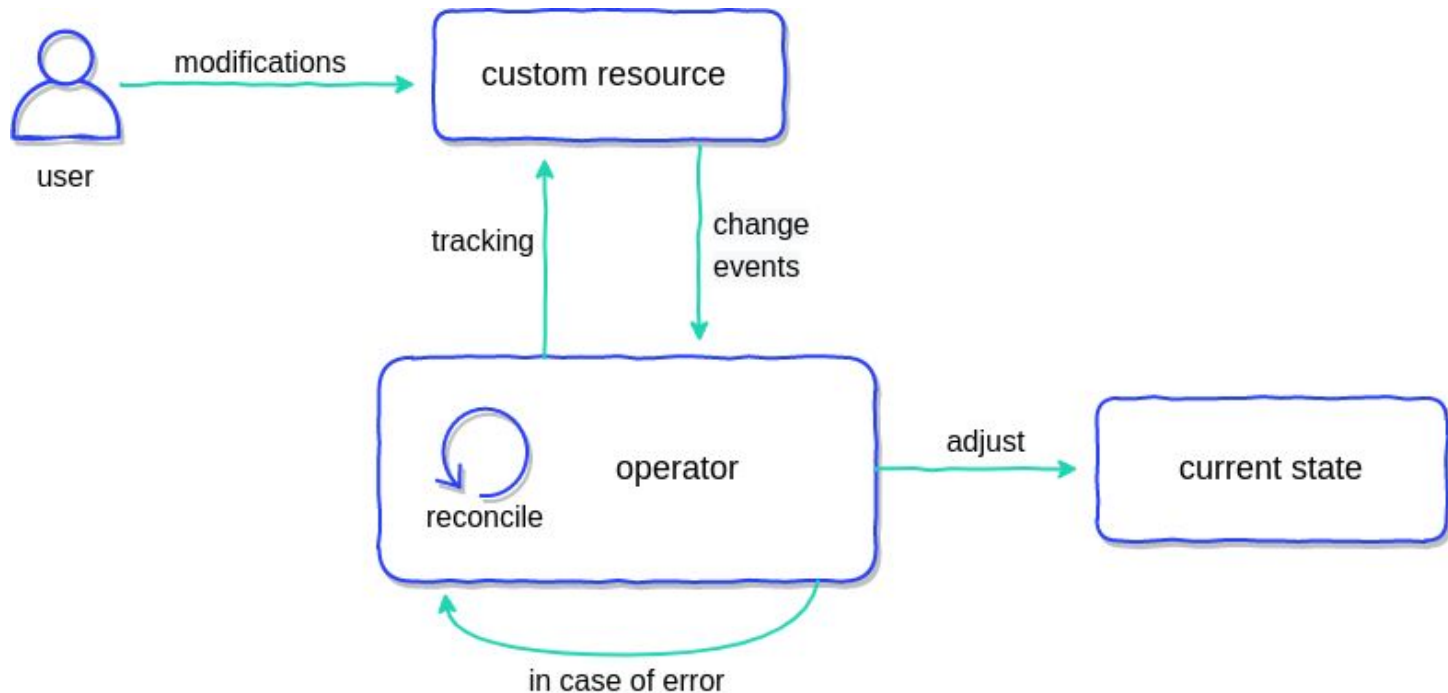


**: A Kubernetes recursion**

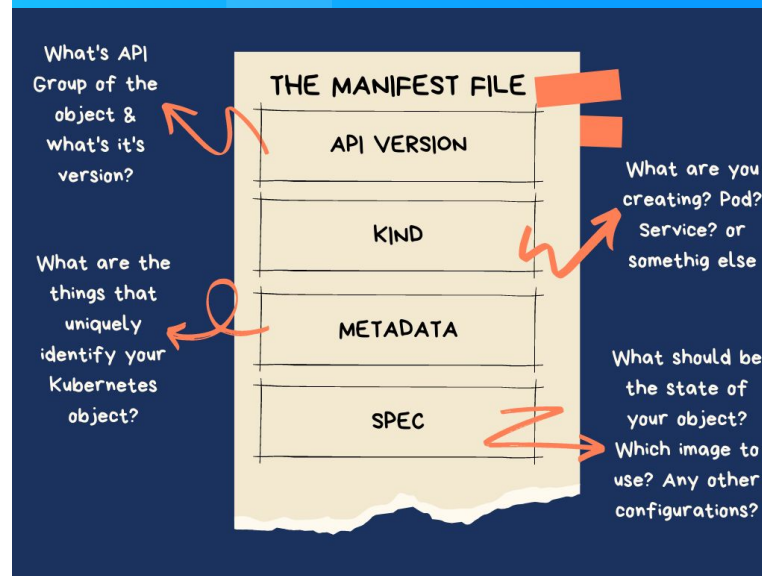
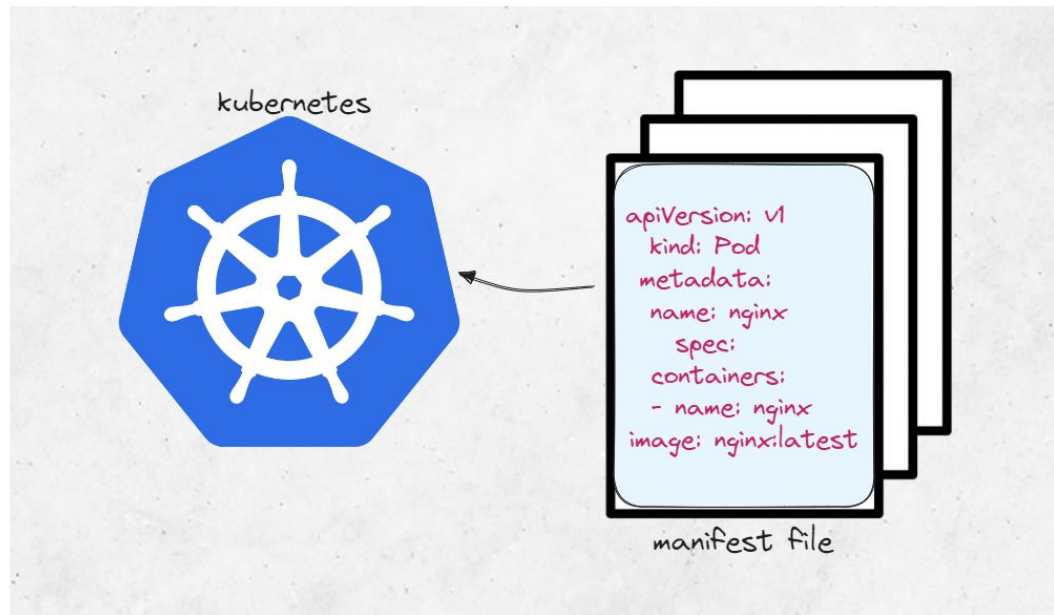
Kubernetes-like control planes for form-factors and use-cases beyond Kubernetes and container workloads.

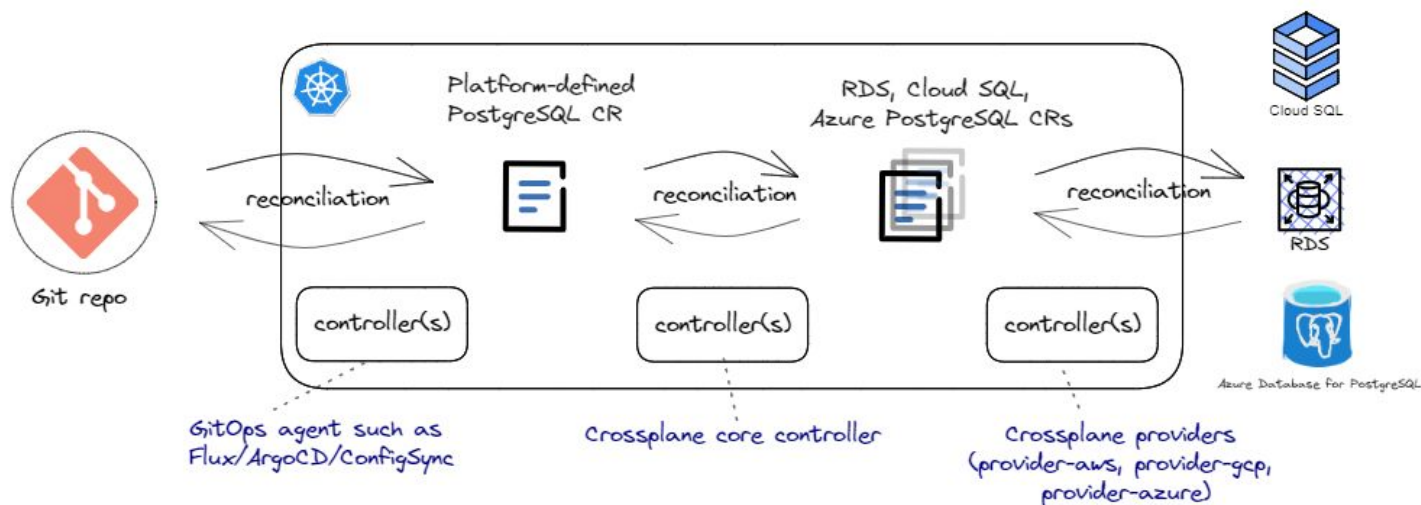
Borislava Bagaliyska, R&D engineer @  
VMware Tanzu Platform

# What makes Kubernetes an awesome platform?



# Manifest files as APIs





Example: GitOps agent + platform defined abstraction + cloud resource provider

- 1st loop: fetches custom resource yaml files that model a custom platform-defined PostgreSQL resource from git repo and syncs them to the API server
- 2nd loop: implements any transformations from custom PostgreSQL resource down to the cloud provider specific PostgreSQL resource definitions
- 3rd loop: talks to cloud APIs and provisions the actual PostgreSQL managed service that the respective cloud offers



“Kubernetes has demonstrated the power of a well architected control plane with a great API.

The industry is beginning to notice that this control plane can be used to do much more than orchestrate containers, and are increasingly looking to use the Kubernetes control plane to manage all of their infrastructure.”

- Crossplane




**I'm a service provider  
and love this  
architecture!  
Can't I just use  
Kubernetes as my  
underlying platform?**




# Not really.



Multi-tenancy




Scalability




It's still a  
container  
platform!




# Enter KCP!



Multi-tenant



Scalable



No pods!

*... but still all the Kubernetes*







## KCP's goals

- Cloud-native control plane for Kubernetes style APIs
- Infrastructure for platform builders
- Full isolation and security of tenants
- Easily extensible
- Bring SaaS to Kubernetes
- Scalable, distributed and fault-tolerant




A man with glasses and a beard, wearing a blue button-down shirt, is seated at a wooden desk in a bright, modern office. He is looking at two computer monitors. The left monitor displays a dashboard with various charts and graphs, while the right monitor shows a web application with a blue header and orange accents. His hands are on a laptop keyboard. The desk is cluttered with various items: a mouse, a small potted plant, a tablet, a notebook, and some cables. In the background, three other people are seated at a long table near a large window, working. The office has a high ceiling with exposed wooden beams and a white radiator.

**Show me how it does it!**



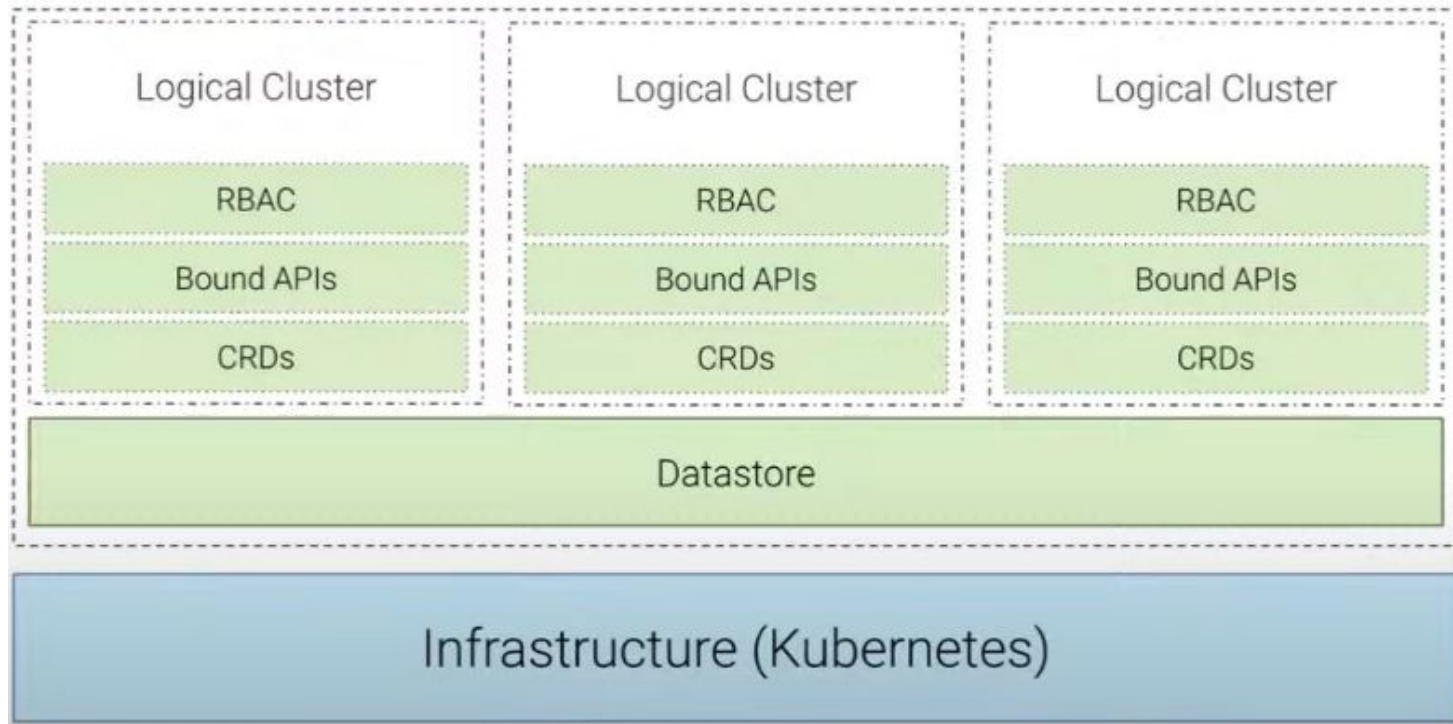
# Tenancy options in Kubernetes



- Namespaces
- Physical clusters
- Clusters in VMs
- vclusters - multiple apiservers running in namespaces of host cluster, workloads reusing host nodes
- HyperShift - same but with separate node pools



# KCP: Logical clusters sharing apiserver & storage



# Logical clusters objects in etcd

(group, version, resource, optional namespace, name)



(group, version, resource, **logical cluster name**, optional namespace, name)




# KCP's tenancy unit - workspaces

`https://kcp:6443/clusters/root`

`https://kcp:6443/clusters/root:org-a`

`https://kcp:6443/clusters/root:org-a:team-a`



# Be your cluster's admin!



# Workspaces are created and navigated with kubectl CLI

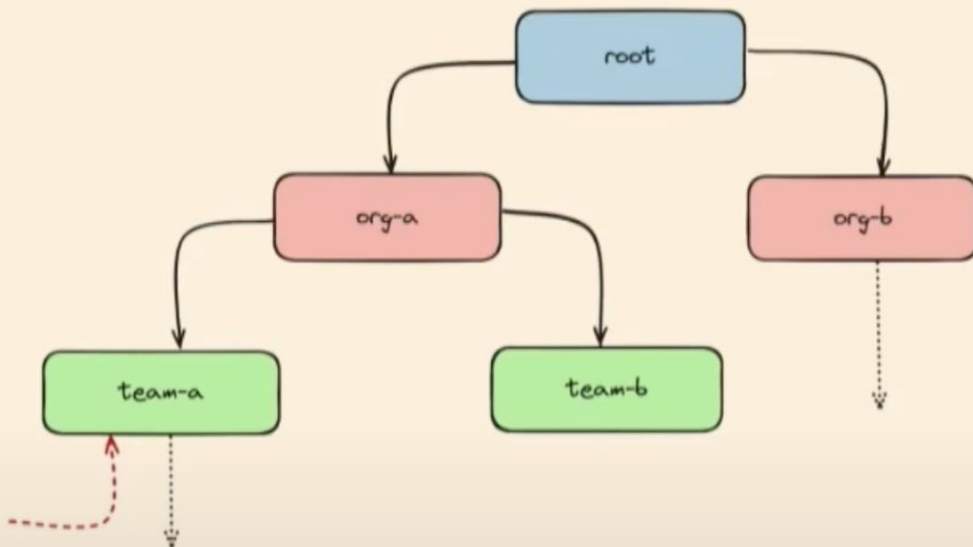
```
$ kubectl ws .  
Current workspace is "root".
```

```
$ kubectl get ws  
NAME      TYPE           REGION  PHASE  URL           AGE  
org-a     organization    Ready   https://...  69d  
org-b     organization    Ready   https://...  65d
```

```
$ kubectl ws org-a  
Current workspace is "root:org-a" (type root:organization).
```

```
$ kubectl get ws  
NAME      TYPE  REGION  PHASE  URL           AGE  
team-a    team   Ready   https://...  3m23s  
team-b    team   Ready   https://...  3m18s
```

```
$ kubectl ws team-a  
Current workspace is "root:org-a:team-a" (type root:team).
```





Try to spot the differences!





# Clusters without workload related resources

<https://kcp:6443/clusters/root:org-a/api>

```
$ kubectl api-resources
```


NAME	SHORTNAMES	APIVERSION	NAMESPACED	KIND
configmaps	cm	v1	true	ConfigMap
events	ev	v1	true	Event
namespaces	ns	v1	false	Namespace
resourcequotas	quota	v1	true	ResourceQuota
secrets		v1	true	Secret
serviceaccounts	sa	v1	true	ServiceAccount
[...]				
<b>workspaces</b>	<b>ws</b>	<b>tenancy.kcp.io/v1alpha1</b>	<b>false</b>	<b>Workspace</b>
workspacetypes		tenancy.kcp.io/v1alpha1	false	WorkspaceType
[...]				




# Adding APIs in workspaces



APIResourceSchema



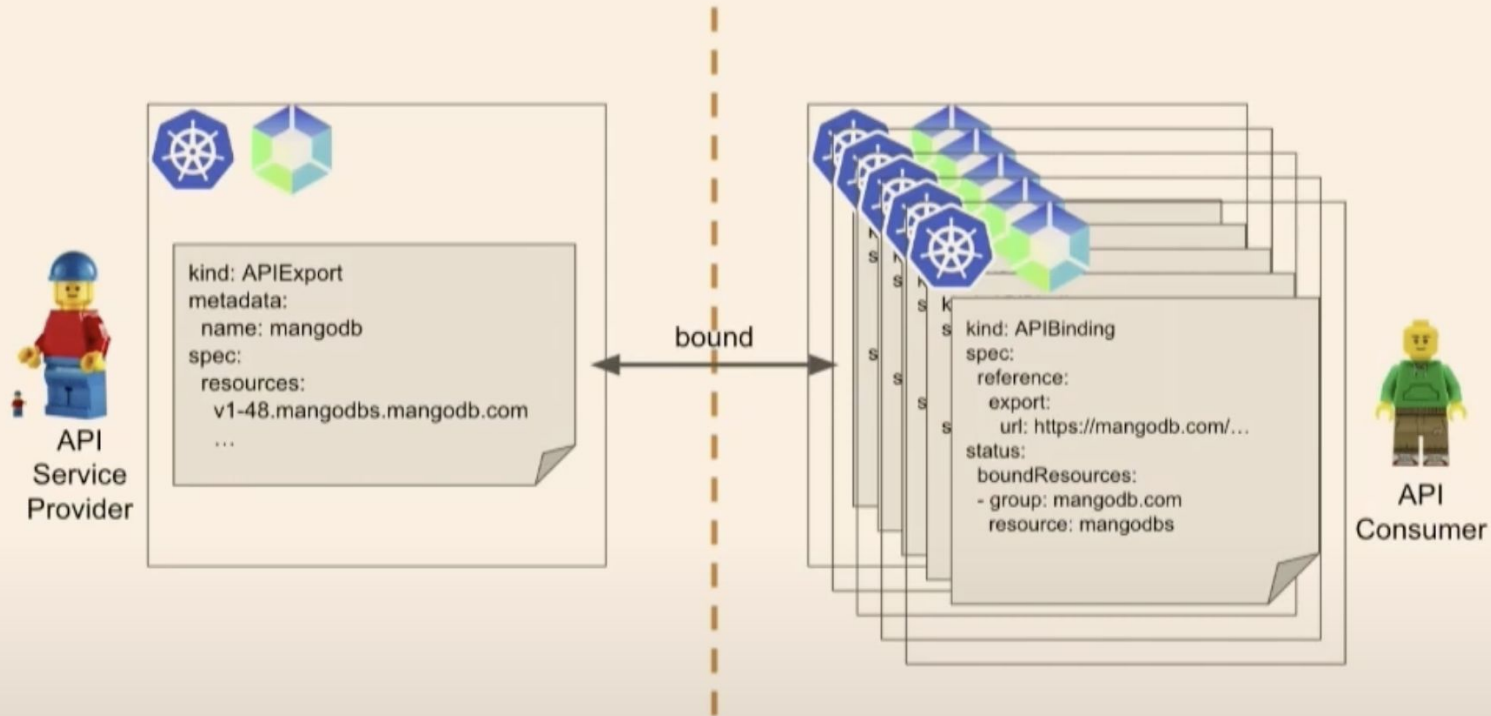
APIExport




APIBinding

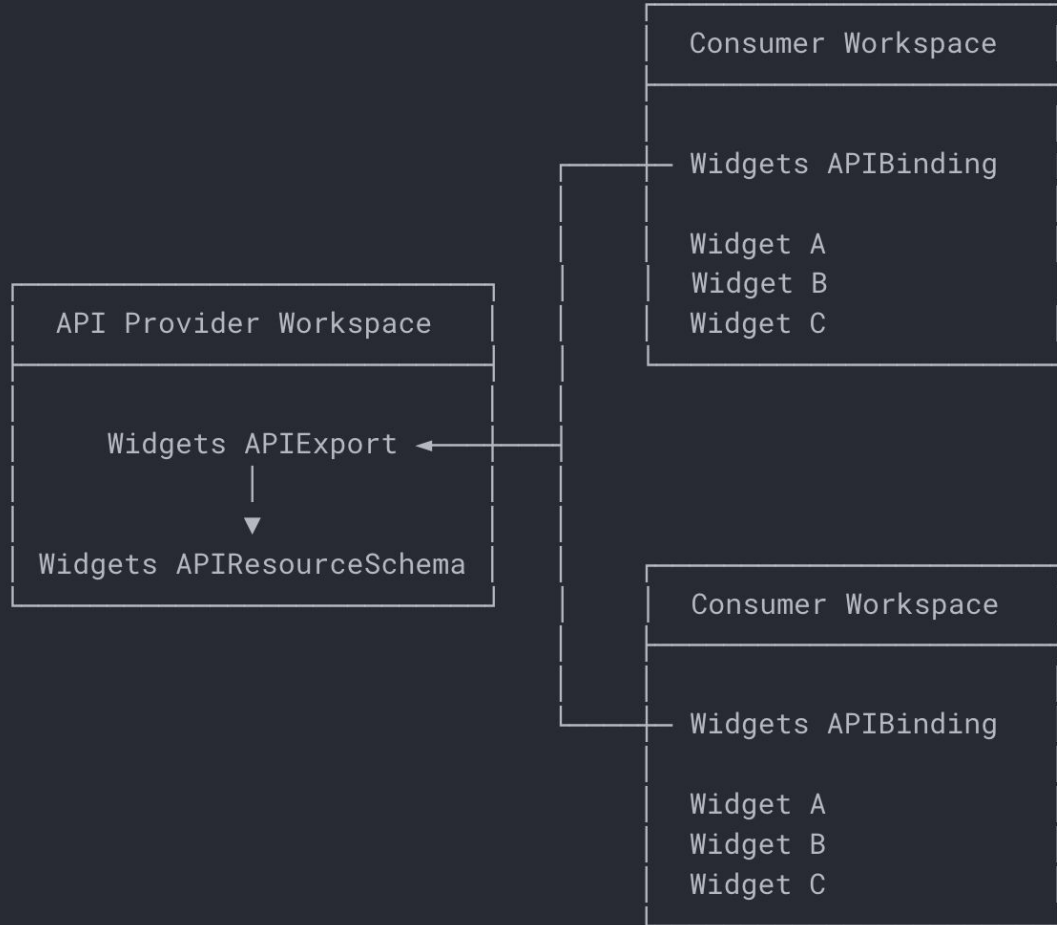


# API bindings



# API bindings





## Consumer Workspace

```
apiVersion: apis.kcp.dev/v1alpha1
kind: APIBinding
spec:
  reference:
    workspace:
      path: root:cert-manager
      exportName: stable
```

# APIBinding

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
spec:
  names:
    kind: Certificate
    listKind: CertificateList
    plural: certificates
    singular: certificate
  ...
```

System shadow workspace

## Service Provider

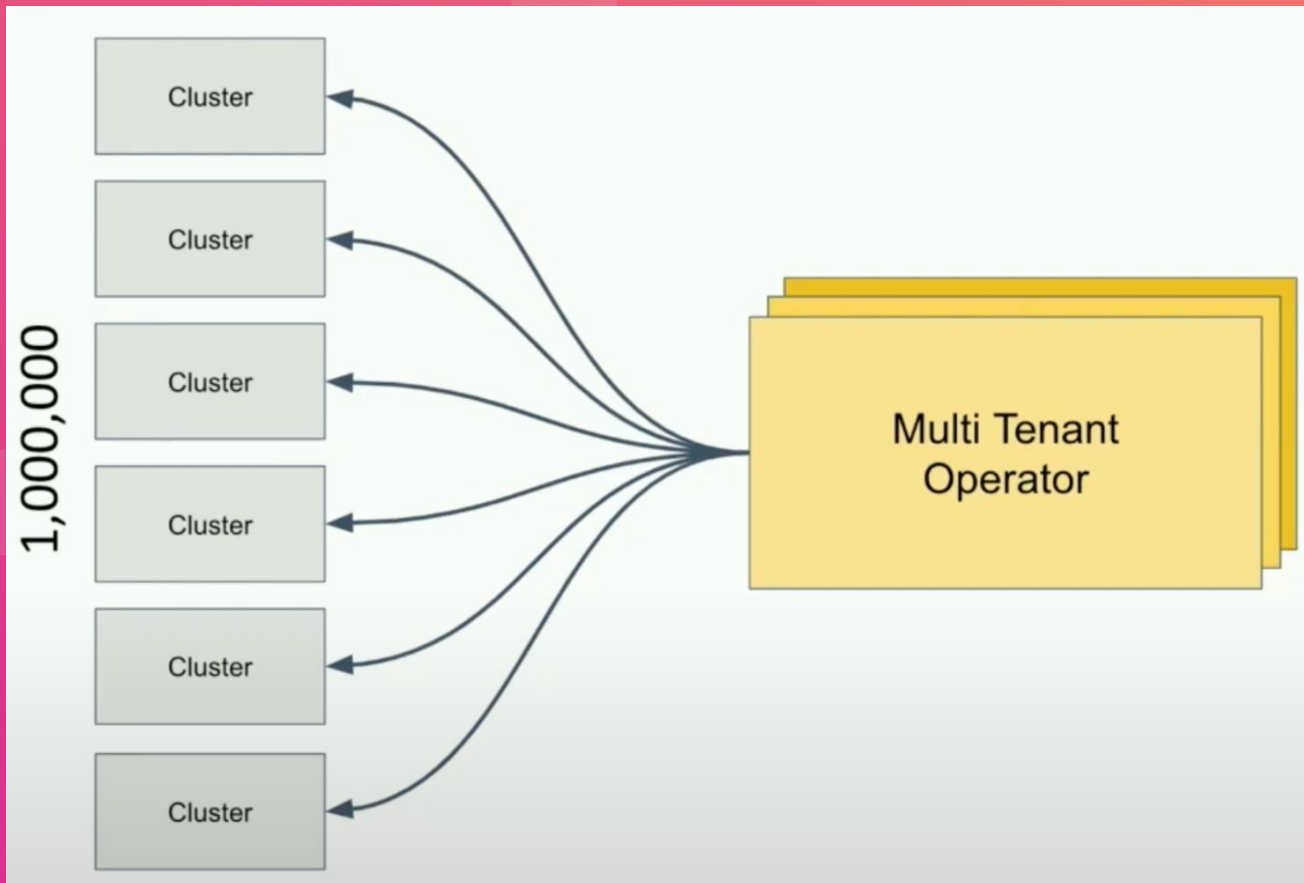
```
apiVersion: apis.kcp.dev/v1alpha1
kind: APIExport
metadata:
  name: stable
spec:
  latestResourceSchemas:
    - certificates.cert-manager.io
    - certificaterequests.cert-manager.io
    - issuers.cert-manager.io
    - clusterissuers.cert-manager.io
```

```
apiVersion: apis.kcp.dev/v1alpha1
kind: APIResourceSchema
spec:
  names:
    kind: Certificate
    listKind: CertificateList
    plural: certificates
    singular: certificate
  ...
```


# APIExport

n:1

# How to reconcile objects in workspaces?

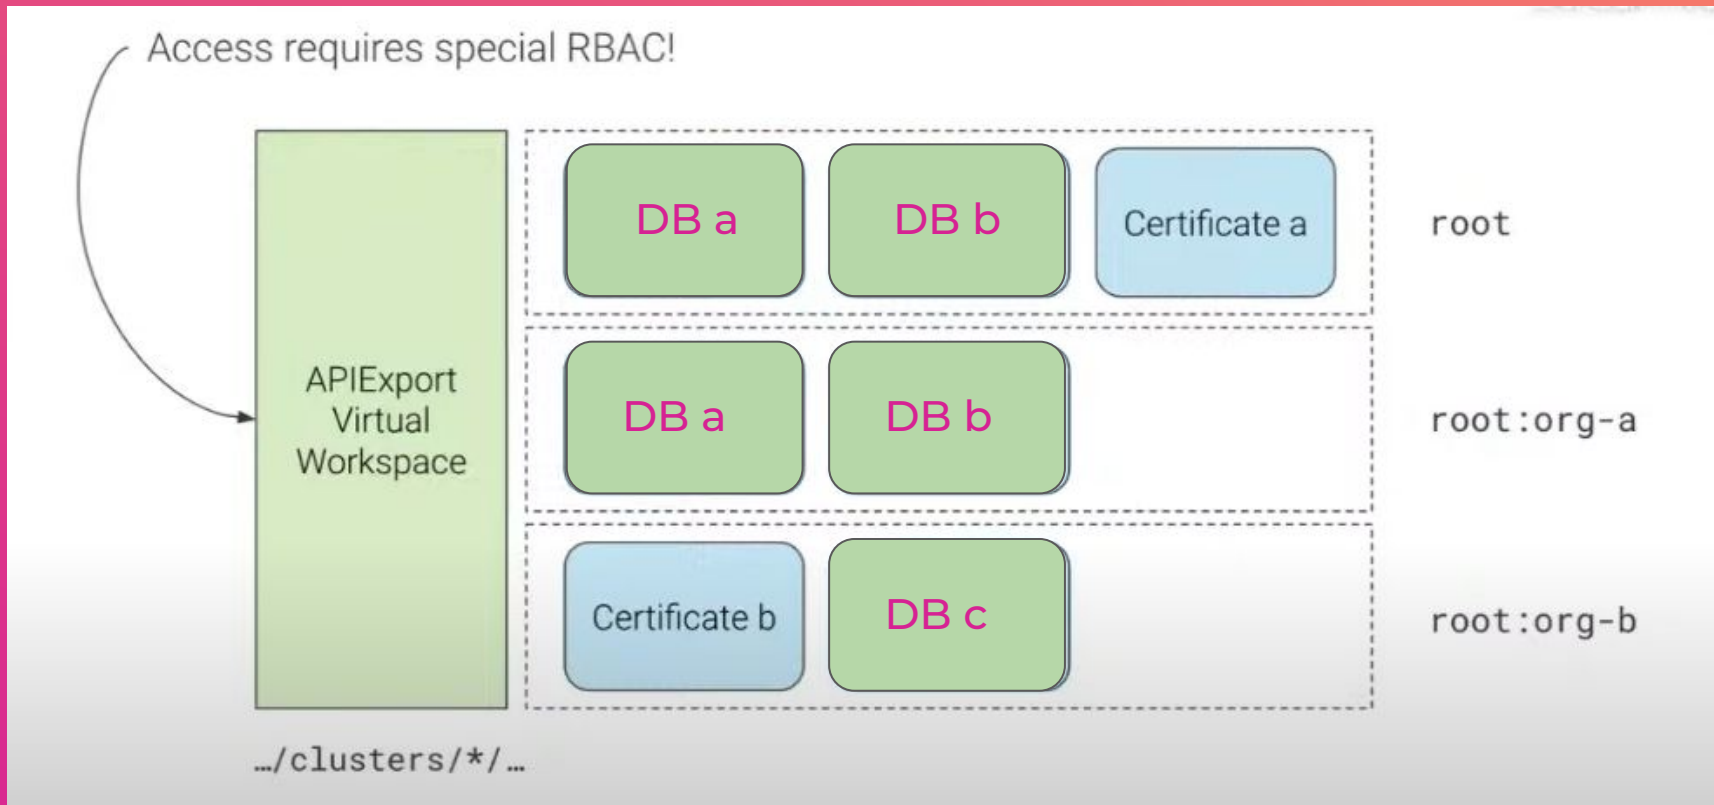


# How does the controller work?





# How does the controller work?





## How is this secure?

1. Organizational access authorizer
2. “**access**” verb for granting access to workspace content
3. “**use**” verb for creating workspaces of particular type
4. Workspace level k8s RBAC
5. “**bind**” verb for creating APIBinding to a particular APIExport
6. Permission claims for granting APIExport controller access to workspace resources
7. Maximal permission policies for exported resources
8. Identity hashes of APIExports



# Requesting and granting permission

A service provider wanting to access `ConfigMaps` needs to specify such a claim in the `APIExport`:

```
spec:
  ...
  permissionClaims:
    - group: ""
      resource: "configmaps"
```

Users can then authorize access to this resource type in their workspace by accepting the claim in the `APIBinding`:

```
spec:
  ...
  permissionClaims:
    - group: ""
      resource: "configmaps"
      state: Accepted
```



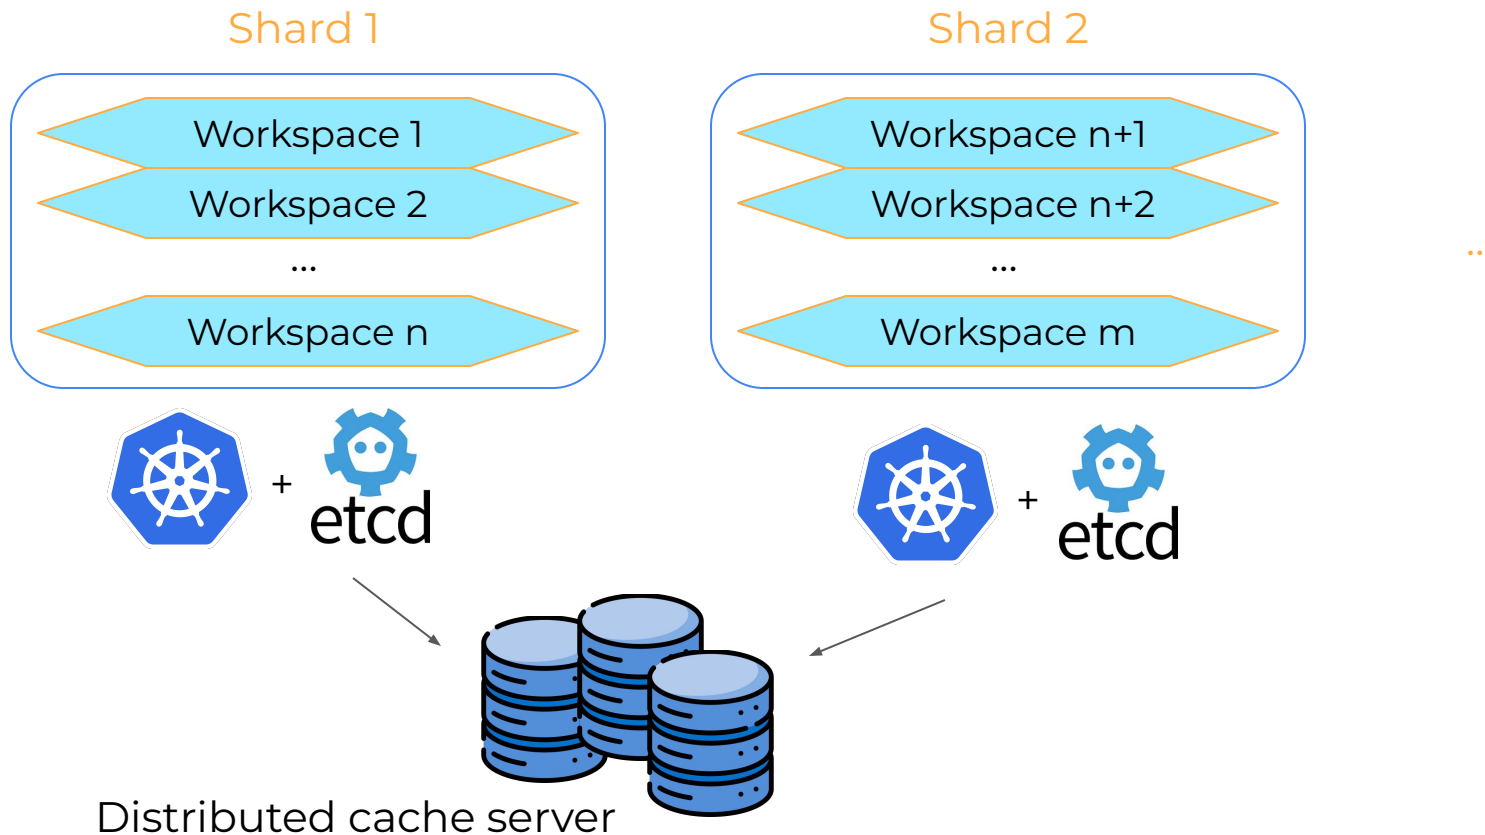
# Advanced permission claims

```
apiVersion: apis.kcp.io/v1alpha1
kind: APIExport
metadata:
  name: example.kcp.io
spec:
  latestResourceSchemas:
  - v220801.widgets.example.kcp.io
  permissionClaims:
  - group: "" # +
    resource: configmaps
    resourceSelector: # +
    - namespace: example-system
      name: my-setup
  - group: somegroup.kcp.io
    resource: things
  identityHash: 5fdf7c7aaf407fd1594566869803f565bb84d22156cef5c445d2ee13ac2cfca6
  all: true # +
```

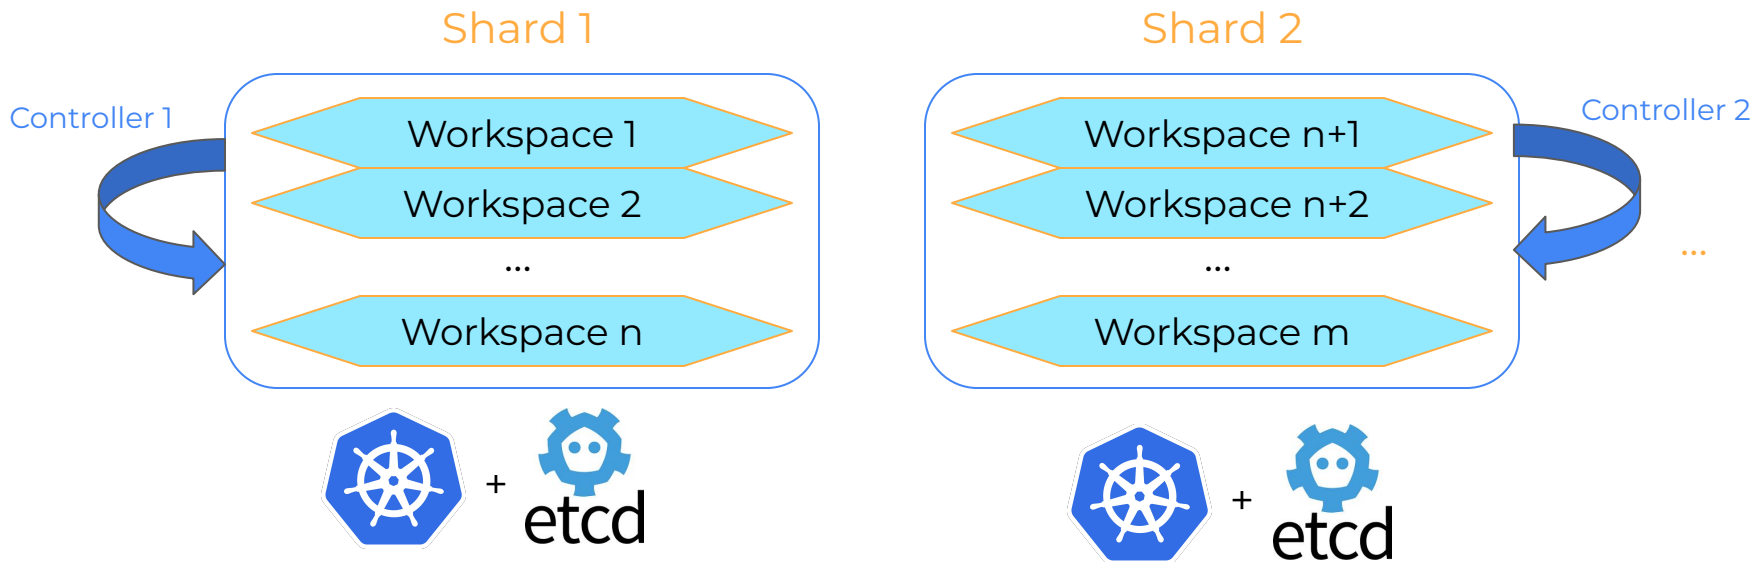


**Scalability by sharding**

# Shards architecture



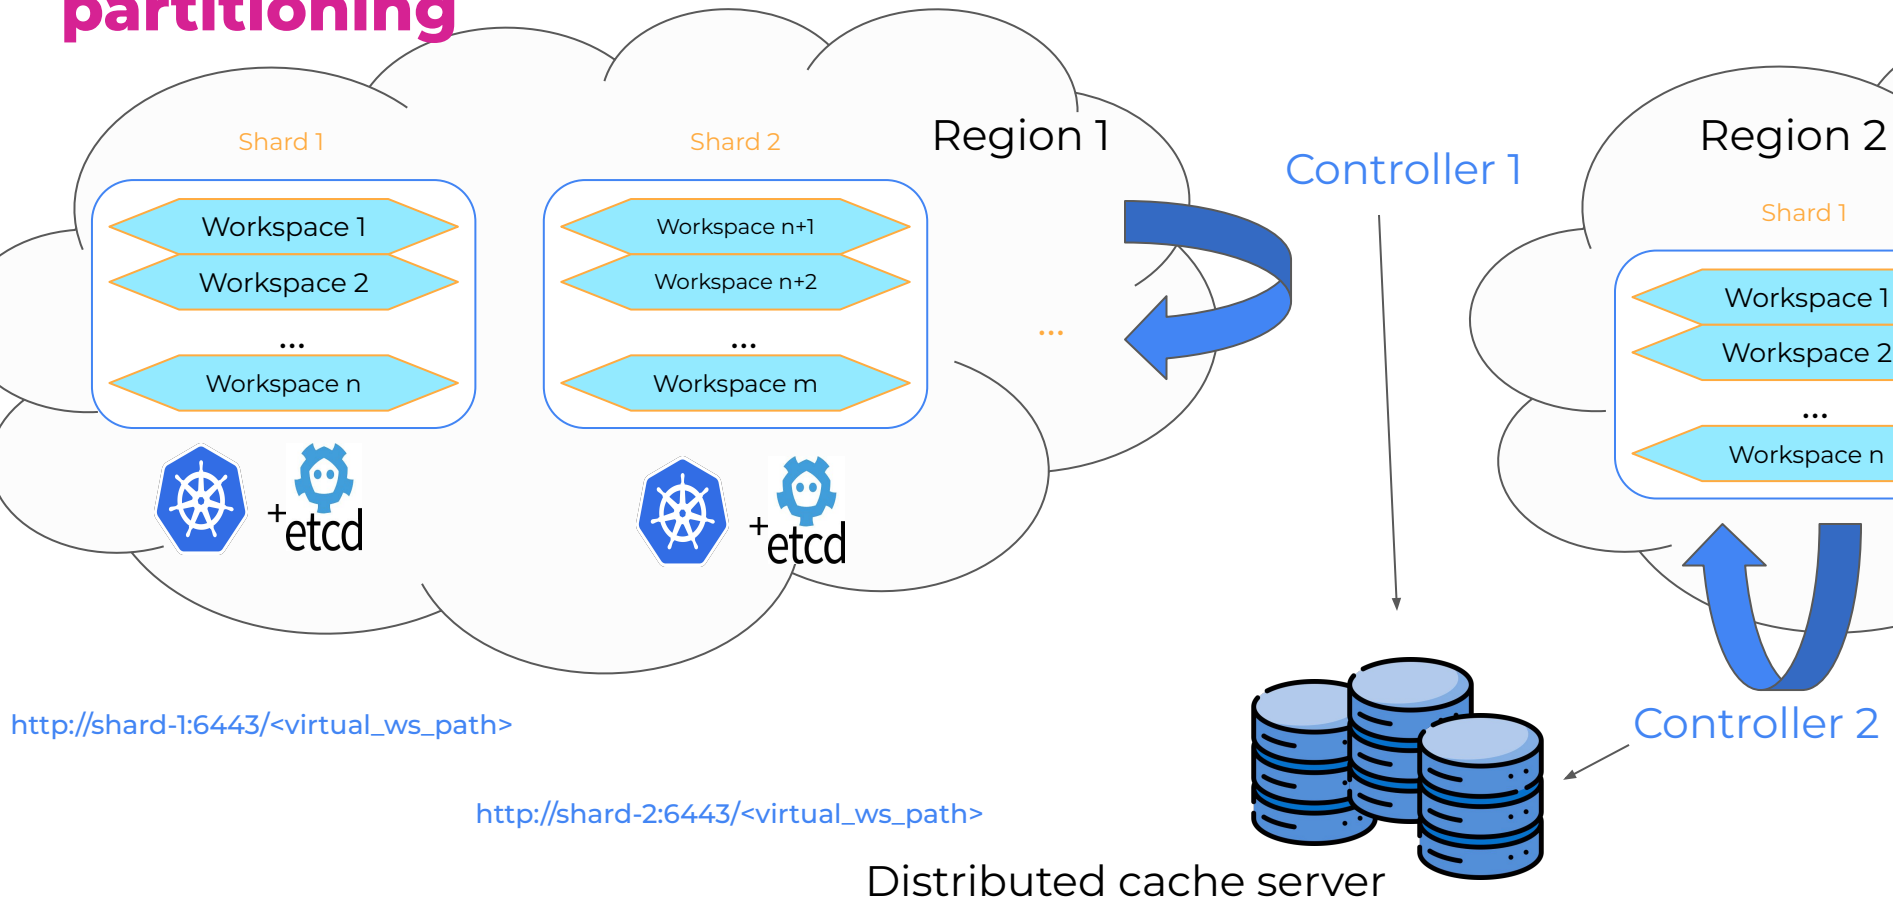
# APIExport controllers in shards



[http://shard-1:6443/<virtual\\_ws\\_path>](http://shard-1:6443/<virtual_ws_path>)

[http://shard-2:6443/<virtual\\_ws\\_path>](http://shard-2:6443/<virtual_ws_path>)

# APIExport controllers in shards - partitioning





# Controller implementation sneak peek

Partitions and APIExportEndpointSlices

```
$ kubectl get apiexportendpointslice example-gcp-europe -o yaml
kind: APIExportEndpointSlice
apiVersion: apis.kcp.io/v1alpha1
metadata:
  name: example-gcp-europe
...
status:
  endpoints
    - url: https://host1:6443/services/apiexport/root/example.kcp.io
    - url: https://host2:6443/services/apiexport/root/example.kcp.io
...
```

## Use cases

**Service  
catalogs**

**IoT & Edge**

**Internal  
developer  
platforms**

**IaaS  
platforms**

e.g. Crossplane,  
Kubernetes LCM

**Transparent  
multi-cluster**



# Shortcomings

## Low maturity

CNCF sandbox project since Sept 2023

## Not suitable for large data

## Eventually consistent by nature





Any questions? 😊



# Resources

- Official docs: <https://docs.kcp.io/kcp/main/>
- KCP repo: <https://github.com/kcp-dev/kcp>
- KCP controller-runtime fork:  
<https://github.com/kcp-dev/controller-runtime>
- Kcp: Towards 1,000,000 Clusters - Stefan Schimanski, Red Hat -  
<https://www.youtube.com/watch?v=fGv5dpQ8X5I>





**CLOUD NATIVE**  
COMPUTING FOUNDATION

**Thank you!**